3.3. Preventing catastrophic forgetting: Experience replay

3.3.1. Catastrophic forgetting

3.3.2. Experience replay

3.4. Improving stability with a target network

3.4.1. Learning instability

3.5. Review

Summary

Chapter 4. Learning to pick the best policy: Policy gradient methods

4.1. Policy function using neural networks

4.1.1. Neural network as the policy function

4.1.2. Stochastic policy gradient

4.1.3. Exploration

4.2. Reinforcing good actions: The policy gradient algorithm

4.2.1. Defining an objective

4.2.2. Action reinforcement

4.2.3. Log probability

4.2.4. Credit assignment

4.3. Working with OpenAI Gym

4.3.1. CartPole

4.3.2. The OpenAI Gym API

4.4. The REINFORCE algorithm

4.4.1. Creating the policy network

4.4.2. Having the agent interact with the environment