8.3. Setting up Super Mario Bros.
8.4. Preprocessing and the Q-network
8.5. Setting up the Q-network and policy function
8.6. Intrinsic curiosity module
8.7. Alternative intrinsic reward mechanisms
Summary

Chapter 9. Multi-agent reinforcement learning

9.1. From one to many agents
9.2. Neighborhood Q-learning
9.3. The 1D Ising model
9.4. Mean field Q-learning and the 2D Ising model
9.5. Mixed cooperative-competitive games
Summary

Chapter 10. Interpretable reinforcement learning: Attention and relational models

10.1. Machine learning interpretability with attention and relational biases
10.1.1. Invariance and equivariance
10.2. Relational reasoning with attention
10.2.1. Attention models
10.2.2. Relational reasoning
10.2.3. Self-attention models
10.3. Implementing self-attention for MNIST
10.3.1. Transformed MNIST
10.3.2. The relational module