- The *return* is the sum of discounted rewards in an episode of the environment.
- A *policy gradient method* is a reinforcement learning approach that tries to directly learn a policy by generally using a parameterized function as a policy function (e.g., a neural network) and training it to increase the probability of actions based on the observed rewards.
- *REINFORCE* is the simplest implementation of a policy gradient method; it essentially maximizes the probability of an action times the observed reward after taking that action, such that each action's probability (given a state) is adjusted according to the size of the observed reward.