- The encoder encodes high-dimensional states into a low-dimensional vector with high-level features (which removes noise and trivial features).
- The forward-prediction model predicts the next encoded state, and its error provides the curiosity signal.
- The inverse model trains the encoder by taking two successive encoded states and predicting the action that was taken.
- Empowerment is a closely related but alternative approach to curiosity-based learning. In empowerment, the agent is incentivized to learn how to maximize the amount of control it has over the environment.