Listing 8.1. Setting up the Super Mario Bros. environment
Listing 8.2. Downsample state and convert to grayscale
Listing 8.3. Preparing the states
Listing 8.4. The policy function
Listing 8.5. Experience replay
Listing 8.6. ICM components
Listing 8.7. Deep Q-network
Listing 8.8. Hyperparameters and model instantiation
Listing 8.9. The loss function and reset environment
Listing 8.10. The ICM prediction error calculation
Listing 8.11. Mini-batch training using experience replay
Listing 8.12. The training loop
Listing 8.13. Testing the trained agent

Chapter 9. Multi-agent reinforcement learning

Listing 9.1. Pseudocode for neighborhood Q-learning, part 1
Listing 9.2. Pseudocode for neighborhood Q-learning, part 2
Listing 9.3. 1D Ising model: Create the grid and produce rewards
Listing 9.4. The 1D Ising model: Generate neural network parameters
Listing 9.5. The 1D Ising model: Defining the Q function
Listing 9.6. The 1D Ising model: Get the state of the environment
Listing 9.7. The 1D Ising model: Initialize the grid
Listing 9.8. The 1D Ising model: The training loop
Listing 9.9. Mean field Q-learning: The policy function
Listing 9.10. Mean field Q-learning: Coordinate and reward functions
Listing 9.11. Mean field Q-learning: Calculate the mean action vector
Listing 9.12. Mean field Q-learning: The main training loop
Listing 9.13. Creating the MAgent environment